

Devops Exercises

Please pick just 2 of the exercises below and complete them. Once you have completed your submission, place it in a github repo that you share.

We will schedule a time to review your solution with you after we have had a chance to review it internally.

AWS VPC Exercise

You have been assigned a project to create a new terraform module for creating a staging VPC in AWS. The following requirements have been given to you.

- The network supernet of 172.16.0.0/16
- The VPC should provide an endpoint for AWS Systems Manager (SSM), internally.
- The VPC should provide an internal endpoint for S3.
- The VPC should have two different availability zones, with two private subnets and two public subnets for each availability zone.
- The VPC should contain appropriate routing tables.
- The VPC should include appropriate NAT gateways and an internet gateway.

Create module code to fulfill this request utilizing terragrunt/terraform. Show how you would use this module code to bring up the given architecture.

What I understood from the exercise related to AWS VPC Exercise

- A **staging VPC** on 172.16.0.0/16
- **Two AZs**, and **two private + two public subnets per AZ** (→ 4 private + 4 public)
- **Routing tables, IGW, NAT gateways** (one per AZ)
- **Private endpoints: SSM (interface) and S3 (gateway)**
- Show **module code** and **how to use it**; optionally via **Terragrunt**; and show **plan only for VPC** resources.

Structure of the files and folder in the mono repository

1. modules/vpc

- Creates the VPC (`cidr = 172.16.0.0/16`) with two AZs.
- 4 private + 4 public subnets (you pass explicit CIDRs).
- **IGW** for public subnets.
- **NAT gateways per AZ** and private route tables with default route via NAT.
- **VPC endpoints submodule:**

- **s3 gateway** endpoint associated to the route tables for internal S3.
- **ssm interface** endpoint (in subnets) with private DNS for internal SSM.

2. modules/s3_bucket_state

- Bootstraps (first run) the **S3 state bucket** and **DynamoDB lock table**.
- Has **creation toggles**: `create_bucket` / `create_lock_table`.
- On later runs we flip them to `false` and feed back the existing names.

3. modules/iam_tf_policies

- Creates/uses IAM policies the pipeline needs:
 - RW to the state bucket + lock table.
 - Minimal VPC “apply” permissions (`Describe` + `CreateVpc`/`CreateTags` etc.).
- Skips creation if an **existing policy ARN** is detected/passed.

4. modules/github_oidc

- Uses an **existing** GitHub OIDC provider ARN (or can create one),
- Creates an **assumable role** for the workflow and **attaches** the policies from (3).

Putting all the code centralized in the mono repository

The root file is `main.tf` calling the modules and passes inputs:

```
module "vpc" {
    endpoints # ← the VPC with subnets, NAT, routes,
    source = "../modules/vpc"
    name = var.vpc_name
    cidr = "172.16.0.0/16"
    azs = ["eu-west-1a", "eu-west-1b"]

    private_subnets =
["172.16.0.0/20", "172.16.16.0/20", "172.16.32.0/20", "172.16.48.0/20"]
    public_subnets =
["172.16.64.0/20", "172.16.80.0/20", "172.16.96.0/20", "172.16.112.0/20"]

    enable_nat_gateway = true
    one_nat_gateway_per_az = true
    single_nat_gateway = false
}

module "s3_bucket_state_oidc" {
    # ← state bucket + lock table (first run only)
    source = "../modules/s3_bucket_state"
    bucket_prefix_name = var.bucket_prefix_name
    lock_table = var.lock_table
    create_bucket = var.create_bucket
    create_lock_table = var.create_lock_table
    state_key = "envs/${var.environment}/terraform.tfstate"
    existing_bucket_name = var.existing_bucket_name
}
```

```

    existing_lock_table = var.existing_lock_table
}

module "iam_tf_policies" {
    # ← policies used by the OIDC role
    source          = "./modules/iam_tf_policies"
    bucket_name     = module.s3_bucket_state_oidc.s3_bucket_id
    lock_table_name = module.s3_bucket_state_oidc.lock_table_name
    region          = var.region
    depends_on      = [module.s3_bucket_state_oidc]
}

module "github_oidc" {
    # ← OIDC role that attaches those policies
    source          = "./modules/github_oidc"
    create_oidc_provider = false
    oidc_provider_arn  = local.existing_provider_arn
    create_oidc_role   = true
    repositories       = var.repository_list
    oidc_role_attach_policies = [
        module.iam_tf_policies.tf_backend_rw_policy_arn,
        module.iam_tf_policies.tf_vpc_apply_policy_arn
    ]
    depends_on = [module.iam_tf_policies]
}

```

Short summary

- The module vpc creates the VPC with subnets, NAT, routes, endpoints.
- The module iam_tf_policies creates the iam tf policy and identify required for github oidc
- The module github_oidc is creating the OIDC role that attaches the policies in the module iam_tf_policies

The dependencies are as follows

- 1) "iam_tf_policies" module depending on "s3_bucket_state_oidc"
- 2) "github_oidc" depending on the "iam_tf_policies" module.

More information about the plan can be seen in [this link](#).

How I split in CI/CD flow

- **Bootstrap job** (long-lived AWS keys just once):
 1. init locally (backend disabled),
 2. create **state S3 + DDB, policies, OIDC role** (module targets),
 3. export bucket/table/role outputs.
- **Deploy job** (assume OIDC role):
 1. flip create_bucket=false and create_lock_table=false,

2. migrate TF backend to S3 + DDB,
3. plan/apply the whole stack.

How I plan the exercise

Three practical ways (pick one):

- **Simple (demo):** target the module

```
terraform plan -target=module.vpc -input=false -no-color
```

(Good for showcasing; do not use `-target` routinely for full lifecycles. It was used for cost reasons)

- **Clean separation:** run the VPC as its **own stack** (Terragrunt)

```
live/staging/vpc/terragrunt.hcl → source = ../../../../modules/vpc
```

Then:

```
cd live/staging/vpc
terragrunt plan
```

- I set inputs so non-VPC modules don't create (e.g. `create_bucket=false`, `create_lock_table=false`) and just `terraform plan` to create VPC resources.

How the requirements for this exercise are met

- CIDR and subnet layout match **172.16.0.0/16** having each **2 AZs × (2 private + 2 public)**.
- **IGW** for public subnets;
- **NAT per AZ** for private egress.
- **Route tables** were created and associated appropriately for public and the private subnets.
- **Endpoints:**
 - **S3: Gateway** endpoint for internal S3 access.
 - **SSM: Interface** endpoint with private DNS for internal SSM.

Deploying An Application

Helm

- Create a simple templated helm chart that loads up an nginx server and an appropriate ingress. (assume either an nginx or alb controller is just fine)
- Create a values.yml file that fulfills the values for this helm chart.
- Show a command that dumps the template generated.

Environment used was vagrant to deploy a local kubernetes cluster by using vagrant up for the files and folder in [this repository](#)

I followed these steps to complete the exercise

1. `helm repo update`

- * This refreshes my locally cached index of all Helm repositories that I have added from my side.
- * It returns the Output shows the repos that were refreshed (`kor`, `ingress-nginx`). You now see the newest chart versions.

2. `helm pull ingress-nginx/ingress-nginx --untar`

This command downloads the ingress-nginx chart and it extracted it into folder `~/ingress-nginx/` so I can edit `values.yaml` and templates locally.

3. I create `metallb-pool.yaml`

4. I launched on the first attempt `kubectl apply -f metallb-pool.yaml`

- * This return the error error: `no matches for kind "IPAddressPool" ... ensure CRDs are installed first`
- * It means that tje MetalLB's **CRDs** (custom resource definitions) and the webhook weren't in the cluster yet,
- * At that stage, Kubernetes service did not detect the `IPAddressPool` and `L2Advertisement` were.

5. Install MetalLB via Helm

```
``bash
kubectl create ns metallb-system
helm repo add metallb https://metallb.github.io/metallb
helm repo update
helm install metallb metallb/metallb -n metallb-system
``
```

- * The command used above was used to install the following services:
 - ** metallb-controller + metallb-speaker (the brains + per-node announcers)
 - ** CRDs and webhook for validating MetalLB resources

6. Immediately applying the pool again → `connect: connection refused` to webhook

- * This is a condition used to validate if the Webhook was created, but it was not ready yet.
- * Then, I ran:

```
``bash
helm upgrade --install metallb metallb/metallb \
-n metallb-system --create-namespace --wait --timeout 3m
``
```

* `--wait` made Helm hold until the controller/webhook were up and serving.

7. Verifying MetalLB readiness

```
```bash
kubectl -n metallb-system get pods
kubectl -n metallb-system get svc metallb-webhook-service
kubectl -n metallb-system get endpoints metallb-webhook-service
kubectl -n metallb-system logs deploy/metallb-controller
```
```

This was the output received

* I saw the controller with teh status Running amd the webhook service were having an endpoint (port 9443)

* The logs showed webhooks enabled which it means that it worked

8. I applied again `metallb-pool.yaml` on my second attempt

```
```bash
metallb-pool.yaml
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
 name: vagrant-pool
 namespace: metallb-system
spec:
 addresses:
 - 192.168.56.240-192.168.56.250

apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
 name: vagrant-l2
 namespace: metallb-system
spec:
 ipAddressPools:
 - vagrant-pool
```
```

9. Now, it worked the `ipaddresspool.metallb.io/vagrant-pool` created` and `l2advertisement` because they were created.

10. I edited `values.yaml` for ingress-nginx to install from the local chart

```
```bash
cd ~/ingress-nginx
helm upgrade --install my-ingress . \
 -n ingress-nginx --create-namespace -f values.yaml --wait
```
```

```
vagrant@controlplane: /ingress-nginx$ helm upgrade --install my-ingress . -n ingress-nginx --create-namespace -f values.yaml --wait
Release "my-ingress" does not exist. Installing it now.
NAME: my-ingress
LAST DEPLOYED: Sun Oct 19 23:25:11 2025
NAMESPACE: ingress-nginx
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The ingress-nginx controller has been installed.
It may take a few minutes for the load balancer IP to be available.
You can watch the status by running 'kubectl get service --namespace ingress-nginx my-ingress-ingress-nginx-controller --output wide --watch'

An example Ingress that makes use of the controller:
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
  namespace: foo
spec:
  ingressClassName: nginx
  rules:
  - host: www.example.com
    http:
      paths:
      - pathType: Prefix
        backend:
          service:
            name: exampleService
            port:
              number: 80
        path: /
# This section is only required if TLS is to be enabled for the Ingress
tls:
- hosts:
- www.example.com
  secretName: example-tls

If TLS is enabled for the Ingress, a Secret containing the certificate and key must also be provided:
apiVersion: v1
kind: Secret
metadata:
  name: example-tls
```

11. I was watching the Service to check if it was running

```
```bash
```

```
kubectl -n ingress-nginx get svc my-ingress-ingress-nginx-controller --watch
```

```
```
```

```
vagrant@controlplane: /ingress-nginx$ kubectl get service --namespace ingress-nginx my-ingress-ingress-nginx-controller --output wide --watch
```

| NAME | TYPE | CLUSTER-IP | EXTERNAL-IP | PORT(S) | AGE | SELECTOR |
|-------------------------------------|--------------|---------------|----------------|----------------------------|-------|---|
| my-ingress-ingress-nginx-controller | LoadBalancer | 172.17.33.105 | 192.168.56.240 | 80:31960/TCP,443:32311/TCP | 4m18s | app.kubernetes.io/component=controller,app.kubernetes.io/instance=my-ingress,app.kubernetes.io/name=ingress-nginx |

12. I saw `TYPE=LoadBalancer` with ****EXTERNAL-IP 192.168.56.240**** — MetalLB assigned the first address from your pool which seems that it worked.

Why those specific errors happened:

- * CRDs were not found.

That is why it required to apply CRs (IPAddressPool/L2Advertisement) before MetalLB installed its CRDs → Kubernetes does not recognize those kinds yet.

- * Webhook connect refused.

The ValidatingWebhookConfiguration existed, but the backing webhook server was not serving (pod still initializing). Using `helm ... --wait` and giving it a moment fixed it.

- * I edited `values.yaml` to ensure that i was in the right path containing the Helm files relative to folder with all helm templates and charts .

- * I switched into the chart folder to apply successfully the manifest files.

13. I verified the final state

```
```bash
```

# MetalLB objects

```
kubectl -n metallb-system get ipaddresspools,l2advertisements
```

```
metallb-system$ kubectl -n metallb-system get ipaddresspools,l2advertisements
vagrant@controlplane:~/ingress-nginx$ kubectl -n metallb-system get ipaddresspools,l2advertisements
NAME AUTO ASSIGN AVOID BUGGY IPS ADDRESSES
ipaddresspool.metallb.io/vagrant-pool true false ["192.168.56.240-192.168.56.250"]

NAME IPADDRESSPOOLS IPADDRESSPOOL SELECTORS INTERFACES
l2advertisement.metallb.io/vagrant-l2 ["vagrant-pool"]
```

# Ingress-NGINX service has an external IP from your pool

```
kubectl -n ingress-nginx get svc my-ingress-ingress-nginx-controller -o wide
```

```
vagrant@controlplane:~/ingress-nginx$ kubectl -n ingress-nginx get svc my-ingress-ingress-nginx-controller -o wide
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE SELECTOR
my-ingress-ingress-nginx-controller LoadBalancer 172.17.33.105 192.168.56.240 80:31960/TCP,443:32311/TCP 35m app.kubernetes.io/component=controller,app.kubernetes.io/instance=ingress-nginx
```

# IngressClass is default (if you set default: true)

```
```bash
```

```
kubectl get ingressclass nginx -o yaml
```

```
```
```

```
vagrant@controlplane:~/ingress-nginx$ kubectl get ingressclass nginx -o yaml | grep -E '^---|^apiVersion|^kind|^metadata|^ annotations|^ meta.helm.sh/release-name|^ meta.helm.sh/release-namespace|^ creationTimestamp|^ generation|^ labels|^ app.kubernetes.io/component|^ app.kubernetes.io/instance|^ app.kubernetes.io/managed-by|^ app.kubernetes.io/name|^ app.kubernetes.io/part-of|^ app.kubernetes.io/version|^ helm.sh/chart|^ name|^ resourceVersion|^ uid|^spec|^ controller'
vagrant@controlplane:~/ingress-nginx$ kubectl get ingressclass nginx -o yaml
apiVersion: networking.k8s.io/v1
kind: IngressClass
metadata:
 annotations:
 meta.helm.sh/release-name: my-ingress
 meta.helm.sh/release-namespace: ingress-nginx
 creationTimestamp: "2025-10-19T23:25:20Z"
 generation: 1
 labels:
 app.kubernetes.io/component: controller
 app.kubernetes.io/instance: my-ingress
 app.kubernetes.io/managed-by: Helm
 app.kubernetes.io/name: ingress-nginx
 app.kubernetes.io/part-of: ingress-nginx
 app.kubernetes.io/version: 1.13.3
 helm.sh/chart: ingress-nginx-4.13.3
 name: nginx
 resourceVersion: "4479"
 uid: f3286580-ae92-4ea0-a257-5fde43a0e4f7
spec:
 controller: k8s.io/ingress-nginx
```

14. Quick smoke test (optional but recommended)

# Deploy a tiny echo app + Service + Ingress:

```
```yaml
```

echo-app.yaml

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: echo
```

```
  namespace: default
```

```
spec:
```

```
  replicas: 1
```

```
  selector:
```

```
    matchLabels: { app: echo }
```

```
  template:
```

```
    metadata:
```

```
      labels: { app: echo }
```

```
    spec:
```



```
containers:
  - name: echo
    image: hashicorp/http-echo
    args: ["-text=hello from ingress"]
    ports:
      - containerPort: 5678
```

```
apiVersion: v1
kind: Service
metadata:
  name: echo
  namespace: default
spec:
  selector: { app: echo }
  ports:
    - port: 80
      targetPort: 5678
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: echo
  namespace: default
spec:
  ingressClassName: nginx
  rules:
    - host: echo.localtest.me # resolves to 127.0.0.1; replace if you prefer
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: echo
                port:
                  number: 80
```

...

```
```bash
kubectl apply -f echo-app.yaml
```
```

15. Test it. Since your LB IP is `192.168.56.240`, either:

* Add a hosts entry: `192.168.56.240 echo.localtest.me`, then:

```
```bash
curl -H "Host: echo.localtest.me" http://192.168.56.240/
```

```
```  
vagrant@controlplane:~/ingress-nginx$ curl -H "Host: echo.localtest.me" http://192.168.56.240/  
hello from ingress  
```
```

\* Change the Ingress host to a name you control and map it to `192.168.56.240` in `/etc/hosts`.

I received `hello from ingress`.

# My `values.yaml` (MetalLB-friendly) recap

```
```yaml  
controller:  
  ingressClassResource:  
    name: nginx  
    default: true  
  service:  
    type: LoadBalancer  
    annotations:  
      metallb.universe.tf/address-pool: vagrant-pool  
      # metallb.universe.tf/loadBalancerIPs: 192.168.56.241 # optional pin  
    externalTrafficPolicy: Local  
  config:  
    proxy-body-size: "64m"  
    enable-brotli: "true"  
```
```

## Summary

The whole flow was the following

- \* Update repos to pull and edit the chart locally
- \* install MetalLB (wait for webhook)
- \* Apply pool/L2Ad to install ingress-nginx with `LoadBalancer`
- \* Get an external IP
- \* Test with a simple app.